

به نام خدا

## **1. Please summarize the following server components:**

### **○ Motherboard (System Board) :**

A motherboard designed for server use generally has two or more CPU sockets typically for Xeon CPUs. Server motherboards also use error-correcting RAM (see ECC) and more PCI Express slots than a desktop computer motherboard. A server motherboard does not require integrated graphics or a sound circuit as server computers are rack mounted without monitors.

### **○ CPU (Central Processing Unit) :**

A central processing unit (CPU), also called a central processor, main processor, or just processor, is the primary processor in a given computer. Its electronic circuitry executes instructions of a computer program, such as arithmetic, logic, controlling, and input/output (I/O) operations. This role contrasts with that of external components, such as main memory and I/O circuitry, and specialized coprocessors such as graphics processing units (GPUs).

The form, design, and implementation of CPUs have changed over time, but their fundamental operation remains almost unchanged. Principal components of a CPU include the arithmetic–logic unit (ALU) that performs arithmetic and logic operations, processor registers that supply operands to the ALU and store the results of ALU operations, and a control unit that orchestrates the fetching (from memory), decoding and execution (of instructions) by directing the coordinated operations of the ALU, registers, and other components. Modern CPUs devote a lot of semiconductor area to caches and instruction-level parallelism to increase performance and to CPU modes to support operating systems and virtualization.

Most modern CPUs are implemented on integrated circuit (IC) microprocessors, with one or more CPUs on a single IC chip. Microprocessor chips with multiple CPUs are called multi-core processors. The individual physical CPUs, called processor cores, can also be multithreaded to support CPU-level multithreading.

An IC that contains a CPU may also contain memory, peripheral interfaces, and other components of a computer; such integrated devices are variously called microcontrollers or systems on a chip (SoC).

- **RAM (Memory)**

Random access memory (RAM) is the hardware in a computing device that provides temporary storage for the operating system (OS), software programs and any other data in current use so they're quickly available to the device's processor. RAM is often referred to as a computer's main memory, as opposed to the processor cache or other memory types.

Random access memory is considered part of a computer's primary memory. It is much faster to read from and write to than secondary storage, such as hard disk drives (HDDs), solid-state drives (SSDs) or optical drives. However, RAM is volatile; it retains data only as long as the computer is on. If power is lost, so is the data. When the computer is rebooted, the OS and other files must be reloaded into RAM, usually from an HDD or SSD.

The term random access, or direct access, as it applies to RAM is based on the facts that any storage location can be accessed directly via its memory address and that the access can be random. RAM is organized and controlled in a way that enables data to be stored and retrieved directly to and from specific locations. Other types of storage -- such as an HDD or CD-ROM -- can also be accessed directly and randomly, but the term random access isn't used to describe them.

Originally, the term random access memory was used to distinguish regular core memory from offline memory. Offline memory typically referred to magnetic tape

from which a specific piece of data could be accessed only by locating the address sequentially, starting at the beginning of the tape.

RAM is similar in concept to a set of boxes organized into columns and rows, with each box holding either a 0 or a 1 (binary). Each box has a unique address that is determined by counting across the columns and down the rows. A set of RAM boxes is called an array, and each box is known as a cell.

To find a specific cell, the RAM controller sends the column and row address down a thin electrical line etched into the chip. Each row and column in a RAM array has its own address line. Any data that's read from the array is returned on a separate data line.

RAM is physically small and stored in microchips. The microchips are gathered into memory modules, which plug into slots in a computer's motherboard. A bus, or a set of electrical paths, is used to connect the motherboard slots to the processor.

RAM is also small in terms of the amount of data it can hold. A typical laptop computer might come with 8 GB or 16 GB of RAM, while a hard disk might hold 10 TB of data. A hard drive stores data on a magnetized surface that looks like a vinyl record. Alternatively, an SSD stores data in memory chips that, unlike RAM, are non-volatile. They don't require constant power and won't lose data if the power is turned off.

### ○ **Storage Drives (HDD/SSD/NVMe)**

Storage drives are hardware components that store and retrieve data in servers and computers. The main types are:

1. HDD (Hard Disk Drive) – Uses spinning magnetic disks; slower but cost-effective for high-capacity storage.
2. SSD (Solid State Drive) – Flash-based, faster and more durable than HDDs, with no moving parts.
3. NVMe SSD – Connects via PCIe for ultra-fast speeds (lower latency than SATA SSDs), ideal for high-performance workloads.

HDDs are best for bulk storage, SSDs balance speed and affordability, while NVMe drives excel in speed-critical applications like databases and virtualization.

### ○ **RAID Controller (Smart Array)**

A RAID controller is also known as a disk array controller, and it is a type of storage component that manages the disk drives in a RAID infrastructure. RAID stands for redundant array of independent disks. A RAID controller is a card or chip located between the operating system and a storage drive (typically a hard drive).

A RAID controller has two primary functions: to combine multiple low speed or low-capacity storage drives into a single faster and higher volume drive so that they work as a logical unit; and to create redundancy to ensure data usability in the event of drive failure.

### ○ **Power Supply Unit (PSU)**

A power supply unit (PSU) converts mains AC to low-voltage regulated DC power for the internal components of a desktop computer. Modern personal computers universally use switched-mode power supplies. Some power supplies have a manual switch for selecting input voltage, while others automatically adapt to the main voltage.

Most modern desktop personal computer power supplies conform to the ATX specification, which includes form factor and voltage tolerances. While an ATX power supply is connected to the mains supply, it always provides a 5-volt standby (5VSB) power so that the standby functions on the computer and certain peripherals are powered. ATX power supplies are turned on and off by a signal from the motherboard. They also provide a signal to the motherboard to indicate when the DC voltages are in spec, so that the computer is able to safely power up and boot. The most recent ATX PSU standard is **version 3.0 as of mid 2024**.

### ○ **Network Interface Card (NIC)**

A network interface controller (NIC, also known as a network interface card, network adapter, LAN adapter and physical network interface) is a computer hardware component that connects a computer to a computer network.

Early network interface controllers were commonly implemented on expansion cards that plugged into a computer bus. The low cost and ubiquity of the Ethernet standard means that most newer computers have a network interface built into the motherboard, or is contained into a USB-connected dongle.

Modern network interface controllers offer advanced features such as interrupt and DMA interfaces to the host processors, support for multiple receive and transmit queues, partitioning into multiple logical interfaces, and on-controller network traffic processing such as the TCP offload engine.

### ○ **Cooling System (Fans and Heat Sinks)**

A cooling system in servers manages heat generated by components like CPUs, GPUs, and power supplies to prevent overheating and ensure stable operation. It typically consists of:

1. Fans – Force air through the chassis to dissipate heat; may be arranged in redundant configurations for reliability.
2. Heat Sinks – Metal (often aluminum or copper) blocks that absorb and spread heat away from hot components like CPUs.
3. Liquid Cooling (in some high-performance servers) – Uses coolant to transfer heat more efficiently than air.

Effective cooling is critical for performance, energy efficiency, and hardware longevity, especially in dense server racks or blade enclosures.

### ○ **Expansion Slots (PCIe)**

An expansion slot is a socket on a computer motherboard that allows you to add additional components to your system. These slots are used to expand the capabilities of your computer and can be used to add new functionality to your system.

There are several different types of expansion slots, each designed to accommodate different types of expansion cards. Some common types of expansion slots include peripheral component interconnect (PCI), PCI express (PCIe), accelerated graphics port (AGP), and industry standard architecture (ISA).

### ○ **Chassis (Rack/Tower/Blade)**

Chassis refers to the physical enclosure that houses servers or computing hardware. Common types include:

1. Rack Chassis – Mounted in server racks (e.g., 1U, 2U, 4U heights), optimized for data centers.
2. Tower Chassis – Standalone, resembling a desktop PC, used for small businesses or workstations.
3. Blade Chassis – Holds multiple blade servers in a modular, high-density design for efficient scaling in data centers.

Each type balances space, cooling, and scalability needs. Rack and blade chassis are common in enterprise environments, while tower chassis suit smaller deployments.

### ○ **BIOS/UEFI Firmware**

BIOS stands for Basic Input/Output System.

It is stored on an EPROM (Erasable Programmable Read-Only Memory), allowing the manufacturer to push out updates easily.

It provides many helper functions that allow reading boot sectors of attached storage and printing things on screen. You can access BIOS during the initial phases of the boot procedure by pressing del, F2 or F10.

UEFI stands for Unified Extensible Firmware Interface. It does the same job as a BIOS, but with one basic difference: it stores all data about initialization and startup in an .efi file, instead of storing it on the firmware.

This .efi file is stored on a special partition called EFI System Partition (ESP) on the hard disk. This ESP partition also contains the bootloader.

### ○ **Backplane**

A backplane or backplane system is a group of electrical connectors in parallel with each other, so that each pin of each connector is linked to the same relative pin of all the other connectors, forming a computer bus. It is used to connect several printed circuit boards together to make up a complete computer system. Backplanes commonly use a printed circuit board, but wire-wrapped backplanes have also been used in minicomputers and high-reliability applications.

A backplane is generally differentiated from a motherboard by the lack of on-board processing and storage elements. A backplane uses plug-in cards for storage and processing.

## **2. What are IPMI and iLO, and what are their functions?**

### **IPMI - Intelligent Platform Management Interface ( General Standard )**

The **Intelligent Platform Management Interface (IPMI)** is a set of computer interface specifications for an autonomous computer subsystem that provides management and monitoring capabilities independently of the host system's CPU, firmware (BIOS or UEFI) and operating system. IPMI defines a set of interfaces used by system administrators for out-of-band management of computer systems and monitoring of their operation. For example, IPMI provides a way to manage a computer that may be powered off or otherwise unresponsive by using a network connection to the hardware rather than to an operating system or login shell. Another use case may be installing a custom operating system remotely. Without IPMI, installing a custom operating system may require an administrator to be physically present near the computer, insert a DVD or a USB flash drive containing the OS installer and complete the installation process using a monitor and a keyboard. Using IPMI, an administrator can mount an ISO image, simulate an installer DVD, and perform the installation remotely.

Using a standardized interface and protocol allows systems-management software based on IPMI to manage multiple, disparate servers. As a message-based, hardware-level interface specification, IPMI operates independently of the operating system (OS) to allow administrators to manage a system remotely in the absence of an operating system or of the system management software. Thus, IPMI functions can work in any of three scenarios:

- before an OS has booted (allowing, for example, the remote monitoring or changing of BIOS settings)
- when the system is powered down
- after OS or system failure – the key characteristic of IPMI compared with in-band system management is that it enables remote login to the operating system using SSH

System administrators can use IPMI messaging to monitor platform status (such as system temperatures, voltages, fans, power supplies and chassis intrusion); to query inventory information; to review hardware logs of out-of-range conditions; or to perform recovery procedures such as issuing requests from a remote console through the same connections e.g. system power-down and rebooting, or configuring watchdog timers. The standard also defines an alerting mechanism for the system to send a simple Network Management Protocol (SNMP) platform event trap (PET).

## **iLO - Integrated Lights-Out ( HP )**

Integrated Lights-Out, or iLO, is a proprietary embedded server management technology by Hewlett Packard Enterprise which provides out-of-band management facilities. The physical connection is an Ethernet port that can be found on most ProLiant servers and microservers of the 300 and above series.

iLO has similar functionality to the lights out management (LOM) technology offered by other vendors, for example, Sun/Oracle's LOM port, Dell DRAC, the IBM Remote Supervisor Adapter and Cisco CIMC.

iLO makes it possible to perform activities on a Compaq, HP or HPE server from a remote location. The iLO card has a separate network connection (and its own IP address) to which one can connect via HTTPS. Possible options are:

- Reset the server (in case the server doesn't respond anymore via the network card)
- Power-up the server (possible to do this from a remote location, even if the server is shut down)
- Remote system console (in some cases however an 'Advanced license' may be required for some of the utilities to work)
- Mount remote physical CD/DVD drive or image (virtual media), depends on license.
- Access the server's Integrated Management Log (IML)
- Can be manipulated remotely through XML-based Remote Insight Board Command Language (RIBCL)
- Full command-line interface support through RS-232 port (shared with system), though the inability to enter function keys prevents certain operations



- SSH remote network access to iLO card supporting public key authentication, 1024 bit DSA key, at least since iLO 3
- iLO Federation
- Two factor authentication
- Remote syslog, depends on license.

### **3. How do IPMI or iLO relate to the BIOS or UEFI firmware?**

#### **1. Separate but Communicate:**

- IPMI/iLO and BIOS/UEFI are separate components.
- The BMC can read and sometimes modify BIOS/UEFI settings remotely.
- This is typically done through a feature called "Virtual KVM" (keyboard, video, mouse) or via specific commands/APIs.

#### **2. Boot-Time Management:**

- Through IPMI or iLO, you can trigger a reboot, enter BIOS/UEFI setup, or change boot order—even if the OS is down.
- The BMC often includes a virtual media feature to mount ISOs remotely and boot from them.

#### **3. Sensor and Event Monitoring:**

- IPMI can access data from the BIOS/UEFI like temperature, fan speed, CPU status, etc.
- BIOS/UEFI may send alerts or system events (like POST errors) to the BMC, which then reports them via IPMI/iLO.

#### **4. Firmware Updates:**

- Some IPMI/iLO interfaces can update the BIOS/UEFI firmware remotely, depending on vendor features and permissions.

### **4. What are CPU sockets on a server, and what is their purpose?**

## 5. Why was the pseudo file system introduced in Linux?

## 6. What are the differences between a pseudo file system and a normal file system?

- Hint: Consider the Linux design philosophy

### What is Pseudo file system?

'Pseudo-' means false, pretend. So "pseudo-filesystem" means a filesystem that doesn't have *actual* files – rather, it has virtual entries that the filesystem itself makes up on the spot.

For example, `/proc` on many OSes is a procfs which dynamically generates directories for every process. Similarly, `/sys` on Linux generates files and directories to represent hardware layouts. There are FUSE-based pseudo-filesystems for a *lot* of things.

`/dev` may be a real filesystem (just a subdirectory of `/`), or a virtual pseudo-filesystem (e.g. devfs), or a middle point such as Linux devtmpfs (which is a full in-memory filesystem but still creates device nodes out of nowhere).

The pseudo filesystem was introduced in Linux to provide a standardized, file-based interface for accessing kernel and system information. Instead of using special tools or syscalls, users and programs can read from or write to files (e.g., in `/proc` or `/sys`) to monitor or configure the system. This simplifies development, supports the Unix philosophy of treating "everything as a file," and makes the system more transparent and flexible.

## 7. What kind of information is available in the `/sys/` directory?

The `/sys` directory in Linux is a vital part of the sysfs virtual filesystem, which provides access to kernel parameters, hardware devices, and system status information.

### 1. `/sys/block`

- This directory contains information and interfaces for block devices, such as hard drives, SSDs, and RAM disks.

- It includes symbolic links to individual block devices (e.g., loop0, ram0, mmcblk0) which represent devices in /devices/.
- It allows interaction with block device parameters like mounting, reading, and writing.

## 2. /sys/bus

- Represents different types of buses in the system, such as PCI, USB, and I2C.
- Contains directories corresponding to the bus type (e.g., pci, usb, i2c).
- For example: The pci directory gives information about PCI devices and their drivers. The usb directory holds details about USB devices connected to the system.

```
RPI:/sys/bus/i2c/devices $ ls
i2c-13 i2c-14 i2c-15

RPI:/sys/bus/spi/devices $ ls
spi10.0
```

## 3. /sys/class

- Groups devices into device classes based on their function, such as block devices, network interfaces, or thermal sensors.
- Key directories include: block: Information on block devices (e.g., hard drives, SSDs). net: Details about network devices like Ethernet or Wi-Fi interfaces. gpio: Provides control over GPIO pins. thermal: Information related to temperature sensors and fans.

```
RPI:/sys/class $ ls | grep i2c
i2c-adapter
i2c-dev

RPI:/sys/class/i2c-dev $ ls
i2c-13 i2c-14 i2c-15
```

```
RPI:/sys/class $ ls | grep spi
spidev
spi_master
spi_slave

RPI:/sys/class/spidev $ ls
spidev10.0
```

#### 4. /sys/devices

- Represents the device tree for the system. It organizes devices based on their drivers and buses.
- This directory shows devices as per their location and classification within the system's architecture.
- For example, `armv8_cortex_a76` might represent a CPU model, or `mmcblk0` could represent an MMC device.

#### 5. /sys/firmware

- Provides information related to the firmware of the system, such as device trees and hardware initialization.
- Notable directories include: `devicetree`: The system's device tree (often used in ARM-based systems). `fdt`: Flattened Device Tree (used for hardware description in embedded systems).

#### 6. /sys/fs

- Contains information about various filesystems and their state.
- Example directories: `bpf`: Information related to BPF (Berkeley Packet Filter). `cgroup`: Control groups for managing resource limits for processes. `ext4`: Information specific to the ext4 filesystem.

#### 7. /sys/kernel

- Contains kernel-related parameters and status.

- Includes important files like: `config`: The kernel configuration used to build the current kernel. `cpu_byteorder`: Indicates the byte order of the CPU (e.g., big-endian or little-endian). `tracing`: Contains data related to kernel tracing, useful for debugging and performance monitoring.

#### 8. `/sys/module`

- Represents kernel modules that are currently loaded into the system.
- Each directory under `/sys/module/` corresponds to a loaded module, and contains configuration, parameters, and state related to that module.

```
RPI:/sys/module $ ls | grep i2c_dev
i2c_dev
RPI:/sys/module $ ls | grep spidev
spidev
```

#### 9. `/sys/power`

- Controls power management features, such as system suspension or hibernation.
- Includes files like `pm_freeze_timeout` and state to manage power states.

#### 10. `/sys/dev`

- Contains information on device nodes that programs can interact with.
- It includes both block and character devices.
- Devices in `/dev` correspond to real hardware devices or virtual devices that user-space programs can access.

#### 11. `/sys/tracing`

- Provides tracing and profiling data, crucial for debugging and performance analysis.

- It allows you to trace kernel functions and system events, which is useful for diagnosing performance bottlenecks or kernel issues.

## **8. What is DMA (Direct Memory Access), and what is its use case in Linux?**

Direct memory access (DMA) is a feature of computer systems that allows certain hardware subsystems to access main system memory independently of the central processing unit (CPU).

Without DMA, when the CPU is using programmed input/output, it is typically fully occupied for the entire duration of the read or write operation, and is thus unavailable to perform other work. With DMA, the CPU first initiates the transfer, then it does other operations while the transfer is in progress, and it finally receives an interrupt from the DMA controller (DMAC) when the operation is done. This feature is useful at any time that the CPU cannot keep up with the rate of data transfer, or when the CPU needs to perform work while waiting for a relatively slow I/O data transfer.

Many hardware systems use DMA, including disk drive controllers, graphics cards, network cards and sound cards. DMA is also used for intra-chip data transfer in some multi-core processors. Computers that have DMA channels can transfer data to and from devices with much less CPU overhead than computers without DMA channels. Similarly, a processing circuitry inside a multi-core processor can transfer data to and from its local memory without occupying its processor time, allowing computation and data transfer to proceed in parallel.

DMA can also be used for "memory to memory" copying or moving of data within memory. DMA can offload expensive memory operations, such as large copies or scatter-gather operations, from the CPU to a dedicated DMA engine. An implementation example is the I/O Acceleration Technology. DMA is of interest in network-on-chip and in-memory computing architectures.

DMA use cases in Linux:

### **1. High-Speed I/O Operations**

- Used by storage devices (SSDs, HDDs) and network cards (NICs) to transfer data efficiently.
- Example: NVMe SSDs use DMA for fast data transfers.

## 2. Graphics Processing

- GPUs use DMA to access framebuffers and textures directly, improving rendering performance.

## 3. Network Packet Handling

- NICs use DMA to transfer packets directly to kernel buffers (e.g., `sk_buff` in Linux), reducing CPU overhead.

## 4. Audio/Video Streaming

- Sound cards and video capture devices use DMA for low-latency data streaming.

## 5. Peripheral Communication (USB, PCIe, SATA)

- DMA is used in buses like PCIe to allow devices to read/write memory without CPU intervention.

## 9. What does the `lsblk` command do internally when executed in Linux?

- **Do `lsusb`, `lspci`, and `lshw` function similarly?**

The `lsblk` command in Linux lists information about all available block devices (e.g., disks, partitions, LVM volumes) in a tree-like format. Internally, it works by gathering data from multiple sources in the Linux kernel and system files. Here's a breakdown of what happens when you run `lsblk`:

### 1. Reading `/sys` Filesystem (`sysfs`)

The primary source of information for `lsblk` is the `/sys` virtual filesystem (`sysfs`), which exposes kernel data structures in a hierarchical manner.

lsblk scans `/sys/block/` to discover all block devices (e.g., `sda`, `nvme0n1`, `vda`).

For each device, it reads attributes like:

- `/sys/block/<device>/size` (device size in sectors)
- `/sys/block/<device>/dev` (major and minor device numbers)
- `/sys/block/<device>/removable` (whether the device is removable)
- `/sys/block/<device>/queue/rotational` (whether it's an SSD or HDD)
- Partition information from `/sys/block/<device>/<partition>/` (e.g., `sda1`).

## 2. Reading `/proc/partitions`

lsblk also checks `/proc/partitions` to get a list of all registered block devices and their major/minor numbers.

## 3. Querying udev (Device Manager)

If available, lsblk may use libudev to gather additional metadata (e.g., device labels, UUIDs, filesystem types).

This involves querying udev databases (typically `/run/udev/data/`) for properties like:

- `ID_FS_TYPE` (filesystem type, e.g., `ext4`, `xfs`)
- `ID_FS_UUID` (filesystem UUID)
- `ID_FS_LABEL` (filesystem label)

## 4. Checking Mount Points

- To determine where devices are mounted, lsblk reads `/proc/mounts` or uses the `stat()` system call on potential mount points.

## 5. LVM and MD (Software RAID) Detection

If Logical Volume Manager (LVM) or Linux Software RAID (`mdadm`) devices exist, lsblk checks:

- `/dev/mapper/` for LVM volumes.
- `/sys/block/md*/` for MD RAID devices.



## 1. Lsusb (List USB Devices)

### Internal Workflow:

- Reads from `/sys/bus/usb/devices/` to discover USB devices.
- Uses `libusb` or directly accesses the USB device filesystem (`/dev/bus/usb/`).
- Queries the kernel's USB subsystem via `usbfs` (USB filesystem).
- Retrieves details like vendor ID, product ID, speed, and device class.
- May use `udev` for additional metadata (e.g., device names).

## 2. Lspci (List PCI Devices)

### Internal Workflow:

- Reads from `/sys/bus/pci/devices/` for PCI device entries.
- Uses `libpci` (via `pciutils` library) to query the PCI configuration space.
- Accesses `/proc/bus/pci/` (legacy) or `/sys/bus/pci/` (modern).
- Retrieves vendor ID, device ID, class, kernel driver in use, and NUMA node (if applicable).
- Can fetch extra details via `-v`, `-vv`, or `-k` (kernel driver info).

## 3. Lshw (List Hardware)

### Internal Workflow:

- Aggregates data from multiple sources:
  - `/proc/cpuinfo` (CPU details)
  - `/proc/meminfo` (RAM)
  - `/sys/class/dmi/id/` (SMBIOS/DMI for motherboard, BIOS, serial numbers)
  - `/sys/bus/` (PCI, USB devices)
  - `/proc/scsi/` (SCSI/SATA devices)
  - `dmidecode` (if available, for firmware-level details)

## 10. How can we simulate a shutdown operation via the /sys file system?

Using /sys/power/state, This file controls system power states (like suspend), not shutdown, but you can simulate similar low-power behavior.

the values freeze, mem, and disk that you can write to /sys/power/state are used to control power-saving sleep states of the Linux system. They don't power off the system but rather suspend it in different ways.

### freeze — Lightweight Suspend (a.k.a. "suspend-to-idle")

```
echo freeze > /sys/power/state
```

- **What it does:** Puts the system in a very low-power idle state.
- **Who does it:** The OS simply idles all CPUs, stops user-space, and turns off unnecessary devices — no hardware coordination.
- **Power saving:** Minimal. CPU is halted but RAM and devices stay powered.
- **Wake-up:** Immediate and very fast.
- **Good for:** Systems where deeper suspend states aren't available.

### mem — Suspend-to-RAM

```
echo mem > /sys/power/state
```

- **What it does:** System state is saved to RAM, CPUs and most devices are powered off.
- **Power saving:** High — only RAM stays powered.
- **Wake-up:** Fast. System resumes from RAM.
- **Risks:** Power loss = data loss (RAM is volatile).
- **Typical usage:** Laptops, phones during suspend/sleep.

### disk — Hibernate (Suspend-to-Disk)

```
echo disk > /sys/power/state
```

- **What it does:** Saves the entire RAM state to disk (usually swap), then powers off the system.
- **Power saving:** Maximum — system is completely powered off.
- **Wake-up:** Slower, because it has to read memory back from disk.
- **Survives reboot or power failure.**

For real shutdown, here's the direct method:

Triggering shutdown via `/proc/sysrq-trigger`

```
echo o > /proc/sysrq-trigger
```

This tells the kernel to power off the system immediately. Here's what the letters mean:

- **o** — shut down and power off the system
- **b** — reboot the system (no unmounting)

## 11. What are the different types of kernels?

- **Monolithic vs. Microkernel vs. Hybrid**
- **What are the advantages and disadvantages of each?**

### 1. Monolithic Kernel

All core functions (e.g., file systems, device drivers, memory management) run in a single large process in kernel space.

Examples: Linux, BSD, early UNIX

Pros:

- Fast performance due to direct function calls
- Efficient hardware communication

Cons:

- Harder to maintain/debug (one bug can crash the whole system)
- Less modular

## 2. Microkernel

Only essential services (like IPC, scheduling) run in kernel space; everything else (drivers, file systems) runs in user space.

Examples: Minix, QNX, L4, seL4

Pros:

- Better stability and security (less kernel code = fewer critical bugs)
- Easier to update/extend components

Cons:

- Slower performance due to user-kernel context switching
- More complex IPC handling

## 3. Hybrid Kernel

A mix of monolithic and microkernel — keeps some services in kernel space for speed, but modular like a microkernel.

Examples: Windows NT, XNU (macOS), modern versions of ReactOS

Pros:

- Good balance of performance and modularity
- Safer than monolithic, faster than microkernel

Cons:

- Still complex and can inherit both monolithic/microkernel drawbacks
- Hard to design properly

## 12. Why is the first sector of a disk used for the MBR?

The first sector of a disk (sector 0) is used for the Master Boot Record (MBR) because it's the first place the BIOS reads during system startup. It contains:

1. Bootloader code (446 bytes) – to load the OS
2. Partition table (64 bytes) – info about disk partitions
3. Boot signature (2 bytes) – marks it as bootable (0x55AA)

**13. If the MBR is located in the first 512 bytes, how does it know the location of GRUB or another bootloader to load the kernel?**

- The MBR (first 512 bytes) contains a small piece of machine code (called stage 1 of the bootloader).
- This code is too small to load the OS directly, so it just knows how to load the next stage of the bootloader (like GRUB stage 1.5 or stage 2).

**How it works step-by-step:**

1. MBR (Stage 1) is loaded into memory by the BIOS.
2. It contains logic to find the next boot code:
  - Either in the "boot sector" of an active partition (via the partition table),
  - Or in the MBR gap (the space between the MBR and the first partition, often ~31 KB), where GRUB stage 1.5 is stored.
3. GRUB stage 1.5 can understand file systems like ext4, and it loads GRUB stage 2 from `/boot/grub/` on the disk.
4. GRUB stage 2 shows the boot menu and loads the kernel.

**14. What are .efi files, and what is their role in the boot process?**

.efi files (Extensible Firmware Interface) are executable files used by modern computers during the boot process, particularly in systems that use UEFI (Unified Extensible Firmware Interface) instead of the older BIOS. These files play a crucial role in loading the operating system and other pre-boot utilities.

**1. UEFI Firmware Execution**

- When a UEFI-based system powers on, the firmware looks for a bootloader (typically an .efi file) stored in the EFI System Partition (ESP).
- The ESP is a FAT32-formatted partition containing essential boot files.

**2. Bootloader (.efi) Execution**

- The main bootloader (e.g., `bootx64.efi` for x64 systems, `bootia32.efi` for 32-bit) is loaded by the UEFI firmware.

**3. Loading the OS Kernel**

- The bootloader reads configuration files (e.g., BCD in Windows, grub.cfg in Linux) to determine how to load the OS kernel.
- It then hands over control to the operating system

#### 4. UEFI Applications & Utilities

- Some .efi files are standalone utilities (e.g., Shell.efi for UEFI Shell, memtest.efi for memory testing).
- These can be launched directly from the UEFI boot menu.

### 15. What is the ESP (EFI System Partition) in UEFI, and how is it used?

The EFI (Extensible Firmware Interface) system partition or ESP is a partition on a data storage device (usually a hard disk drive or solid-state drive) that is used by computers that have the Unified Extensible Firmware Interface (UEFI). When a computer is booted, UEFI firmware loads files stored on the ESP to start operating systems and various utilities.

An ESP contains the boot loaders, boot managers, or kernel images of installed operating systems (which are typically contained in other partitions), device driver files for hardware devices present in a computer and used by the firmware at boot time, system utility programs that are intended to be run before an operating system is booted, and data files such as error logs.

### 16. Please explain the following section from /etc/grub/grub.conf:

```
menuentry 'Ubuntu' --class ubuntu --class gnu-linux --class gnu
--class os
$menuentry_id_option
'gnulinux-simple-3e3d2181-a1f5-4456-867c-a69f52c910e6'
{
recordfail
load_video
gfxmode $linux_gfx_mode
insmod gzio
```

```

if [ x$grub_platform = xen ]; then insmod xzio; insmod lzopio;
fi
insmod part_gpt
insmod ext2
set root='hd0,gpt2'
if [ x$feature_platform_search_hint = xy ]; then
search --no-floppy --fs-uuid --set=root --hint-bios=hd0,gpt2
--hint-efi=hd0,gpt2 --hint-baremetal=ahci0,gpt2
49ee8c4e-7d13-455b-b287-488a33286e30
else
search --no-floppy --fs-uuid --set=root
49ee8c4e-7d13-455b-b287-488a33286e30
fi
linux /vmlinuz-5.4.0-65-generic root=/dev/mapper/vg0-root ro
maybe-ubiquity
initrd /initrd.img-5.4.0-65-generic
}

```

## 1. Menu Entry Definition

```

menuentry 'Ubuntu' --class ubuntu --class gnu-linux --class gnu
--class os

```

- menuentry 'Ubuntu': Creates a boot menu entry labeled "Ubuntu".
- --class: Assigns classes (ubuntu, gnu-linux, gnu, os) for theming and grouping

## 2. Menu Entry ID

```

$menuentry_id_option
'gnulinux-simple-3e3d2181-a1f5-4456-867c-a69f52c910e6'

```

- A unique identifier for this entry (often derived from the OS or disk UUID).

### 3. Boot Configuration

#### Basic Settings

```
recordfail      # Marks if the last boot failed (for recovery options).
load_video      # Loads video drivers for better display.
gfxmode $linux_gfx_mode # Sets the graphics mode (resolution).
```

- Ensures proper display and tracks boot failures.

#### Compression & Filesystem Modules

```
insmod gzio      # Loads GZIP decompression support.
if [ x$grub_platform = xxen ]; then insmod xzio; insmod lzopio;
fi # Xen-specific modules.
insmod part_gpt  # Supports GPT partition tables.
insmod ext2      # Supports EXT2/3/4 filesystems.
```

- Loads necessary modules for disk access and decompression.

### 4. Root Filesystem Detection

```
set root='hd0,gpt2' # Sets root to the 2nd partition on the 1st disk (GPT).
```

- hd0,gpt2 = First disk (hd0), second partition (gpt2).

#### Search for Root by UUID

```
if [ x$feature_platform_search_hint = xy ]; then
    search --no-floppy --fs-uuid --set=root --hint-bios=hd0,gpt2
    --hint-efi=hd0,gpt2 --hint-baremetal=ahci0,gpt2
    49ee8c4e-7d13-455b-b287-488a33286e30
else
```



```
search --no-floppy --fs-uuid --set=root  
49ee8c4e-7d13-455b-b287-488a33286e30  
fi
```

- `search --fs-uuid` locates the partition with UUID 49ee8c4e-7d13-455b-b287-488a33286e30.
- `--hint-*` provides hints for BIOS/EFI/disk controllers to speed up search.

## 5. Kernel & Initramfs Loading

```
linux /vmlinuz-5.4.0-65-generic root=/dev/mapper/vg0-root ro maybe-ubiquity
```

linux: Loads the Linux kernel (vmlinuz-5.4.0-65-generic).

- `root=/dev/mapper/vg0-root`: Specifies the root filesystem (LVM volume).
- `ro`: Mounts root as read-only initially (later remounted `rw`).
- `maybe-ubiquity`: Flag for Ubuntu installer (if present).

```
initrd /initrd.img-5.4.0-65-generic
```

- Loads the initial RAM disk (initrd), which contains drivers and setup scripts for early boot.